

Introduction to the **INLA** package

Elias T. Krainski

Jul-2017, Lavras/MG
62^a RBras & 17^o SEAGRO

Outline

1 Overview

2 Tokyo example

3 GAG urine data

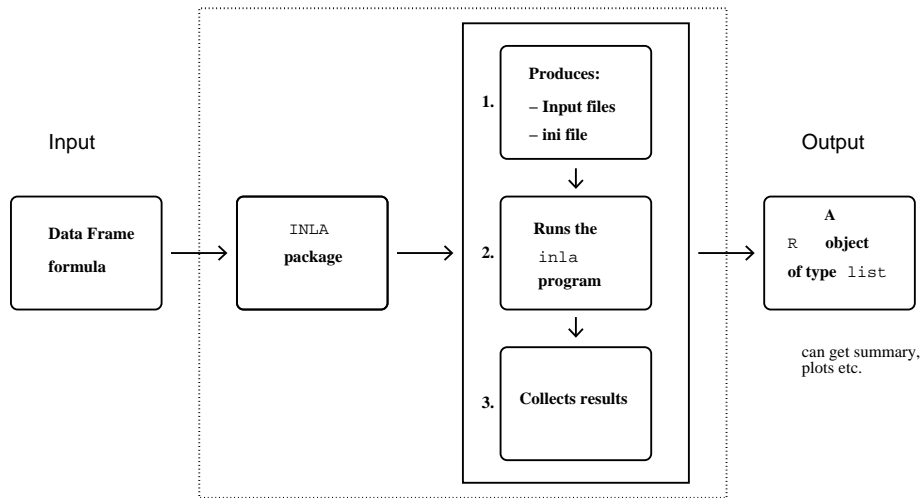
4 Close look at GAG urine example

5 Details

What is **INLA**?

- The algorithm name:
Integrated Nested Laplace Approximations - **INLA**
- **INLA**:
 - The **R**-package that implements **INLA**
 - **INLA** is a collection of R-code and C-code
 - C-code: fast computations for **GMRf**
 - R-code: to have it available for R-users
 - available at www.r-inla.org
- all the source code is public available at
<https://bitbucket.org/hrue/r-inla>

INLA scheme



How to install **INLA**?

new way (from 3rd January 2015)

```
> install.packages("INLA",  
  repos="https://inla.r-inla-download.org/R/stable")
```

or

```
> install.packages("INLA",  
  repos="https://inla.r-inla-download.org/R/testing")
```

INLA is in continuous development:

```
> inla.upgrade()
```

or

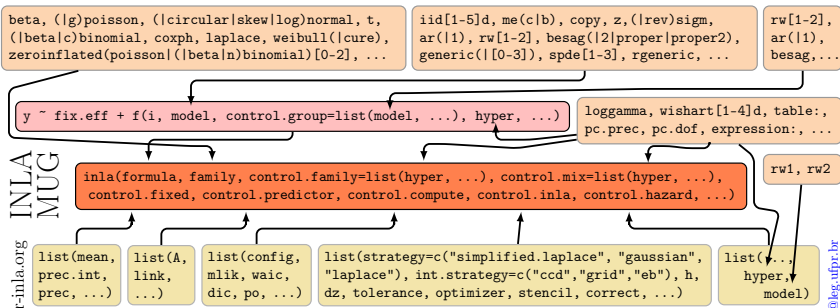
```
> inla.upgrade(testing=TRUE)
```

Data analysis steps using INLA

Five steps to fit an model

- ① data organization
response, covariates and index sets (for random effects)
- ② Define the priors
- ③ define the formula
 - notation (similar to `lm` and `glm` functions)
 - `f(.)` term to define random effects
- ④ call to the `inla`-program
- ⑤ extraction of posterior information

INLA picture



Outline

1 Overview

2 Tokyo example

3 GAG urine data

4 Close look at GAG urine example

5 Details

- Observation model: $p(\mathbf{y}|\mathbf{x})$ likelihood

$$y_i \sim \text{Binomial}(n_i, p_i)$$

$$p_i = \frac{1}{1 + \exp(-x_i)}$$

$$\pi(\mathbf{y}|\mathbf{x}) = \prod_{i=1}^{366} \pi(y_i|x_i)$$

- $\pi(\mathbf{x}|\theta)$

- 1 first order random walk

$$x_{a_i} - x_{a_{i-1}} \sim N(0, \sigma_1^2)$$

- 2 second order random walk (smoother)

$$x_{a_i} - (2x_{a_{i-1}} - x_{a_{i-2}}) \sim N(0, \sigma_2^2)$$

- $\pi(\sigma^2) = \pi(\theta)$: Penalized Complexity (PC) prior

- define a base/simplest model
- $p(x|\psi)$: ψ is a flexibility parameter
- $\psi = 0$ is the base model case

- define a base/simplest model
- $p(x|\psi)$: ψ is a flexibility parameter
- $\psi = 0$ is the base model case
- **Idea:** a prior that has mode at the base model
- posterior concentrates on $\psi > 0$ only if the data requires a more complex model

R-code for Tokyo example rw1

The dataset

```
data(Tokyo)
Tokyo[1:3,]

##      y n time
## 1 0 2     1
## 2 0 2     2
## 3 1 2     3
```

Define the formula (include hyperprior)

```
pcprec <- list(theta=list(
  prior='pc.prec', param=c(2, 0.05)))
formula1 <- y ~ 0 +
  f(time, model='rw1',
    cyclic=TRUE, constr=FALSE,
    scale.model=TRUE, hyper=pcprec)
```

```
tokyo1 <- inla(formula1, data=Tokyo, Ntrials=n,
  family='binomial', control.predictor=list(compute=TRUE),
  control.compute=list(dic=TRUE, cpo=TRUE))
```

R-code: Tokyo example rw2

Define the formula

```
formula2 = y ~ 0 +  
  f(time, model='rw2', cyclic=TRUE, constr=FALSE,  
    scale.model=TRUE, hyper=pcprec)
```

Fit the model

```
tokyo2 <- inla(formula2, data=Tokyo, Ntrials=n,  
  family='binomial', control.predictor=list(compute=TRUE),  
  control.compute=list(dic=TRUE, cpo=TRUE))
```

Tokyo rw1 summary

```
summary(tokyo1)

##
## Call:
## c("inla(formula = formula1, family = \"binomial\", data = Tokyo, Ntrials = n, \"      control.compute = lis
##
## Time used:
##   Pre-processing      Running inla Post-processing      Total
##           0.3272           0.8145           0.0702           1.2119
##
## The model has no fixed effects
##
## Random effects:
##   Name      Model
##   time      RW1 model
##
## Model hyperparameters:
##           mean      sd 0.025quant 0.5quant 0.975quant  mode
## Precision for time 1.049 0.7564      0.2807  0.8384      3.113 0.5629
##
## Expected number of effective parameters(std dev): 22.08(6.605)
## Number of equivalent replicates : 16.57
##
## Deviance Information Criterion (DIC) ...: 616.70
## Effective number of parameters .....: 23.27
##
## Marginal log-Likelihood: -944.22
## CPO and PIT are computed
##
## Posterior marginals for linear predictor and fitted values computed
```

Some Tokyo results rw1

Random effects

```
round(tokyo1$summary.random$time[1:3,], 4)
```

##	ID	mean	sd	0.025quant	0.5quant	0.975quant	mode	kld
## 1	1	-1.6366	0.4452	-2.5249	-1.6345	-0.7625	-1.6317	0
## 2	2	-1.6144	0.4450	-2.4992	-1.6134	-0.7381	-1.6127	0
## 3	3	-1.5759	0.4468	-2.4573	-1.5774	-0.6880	-1.5814	0

Fitted values

```
round(tokyo1$summary.fitted.values[1:3,], 4)
```

##		mean	sd	0.025quant	0.5quant	0.975quant	mode
## fitted.Predictor.001		0.1717	0.0626	0.0742	0.1632	0.3182	0.1473
## fitted.Predictor.002		0.1748	0.0635	0.0760	0.1661	0.3235	0.1499
## fitted.Predictor.003		0.1804	0.0655	0.0790	0.1712	0.3344	0.1541

Compare

Precision

```
round(rbind(rw1=tokyo1$summary.hy[, 1:6],  
            rw2=tokyo2$summary.hy[, 1:6]), 4)
```


```
##          mean      sd 0.025quant 0.5quant 0.975quant      mode  
## rw1  1.0492 0.7564      0.2807   0.8384      3.1128 0.5629  
## rw2  1.8688 3.7269      0.1095   0.6424     11.8073 0.2035
```

DIC, -LCPO

```
round(rbind(DIC=c(rw1=tokyo1$dic$dic, rw2=tokyo2$dic$dic),  
            LCPO=c(sum(-log(tokyo1$cpo$cpo)), sum(-log(tokyo2$cpo$cpo)))))
```

```
##          rw1      rw2  
## DIC    616.7 635.29  
## LCPO    307.9 317.68
```



Tokyo results




Outline

- 1 Overview
- 2 Tokyo example
- 3 GAG urine data
- 4 Close look at GAG urine example
- 5 Details

GAG urine



GAG urine aim



- Problem: build a chart for GAG urine

$$E(GAG|age) = \mu(age)$$

- $\mu(age)$ must be positive
→ **Easy**: model for the log of GAG
- How to relate GAG and age?
→ People usually use parametric non-linear models
- but we have INLA...
→ Let's define a semi-parametric one

GAG urine model

- linear model for $y = \log(\text{GAG})$

$$y = \beta_0 + f(\text{age}) + \text{error}$$

$$\text{error} \sim N(0, \sigma_e^2)$$

- non linear effect from age

$$f(\text{age}) = x_{\text{age}}$$

where x for some age is similar to x for **neighbour** ages

- define some grid of ages: a_1, a_2, \dots, a_k
- first order random walk

$$x_{a_i} - x_{a_{i-1}} \sim N(0, \sigma^2)$$

- second order random walk (smoother)

$$x_{a_i} - (2x_{a_{i-1}} - x_{a_{i-2}}) \sim N(0, \sigma_x^2)$$

GAG urine analysis

GAG urine for 314 children aged from 0 to 17 years

```
require(MASS)
data(GAGurine)
GAGurine[c(1, 314), ]
```

```
##      Age  GAG
## 1    0.00 23.0
## 314 17.67  9.3
```

Prepare the data

```
GAGurine$y =
  log(GAGurine$GAG)
```

Define the formula

```
gag.form = y ~
  f(inla.group(Age), model='rw2',
    scale.model=TRUE, hyper=pcprec)
```

Fit the model

```
gag.rw2 <- inla(gag.form,
  data=GAGurine)
```

1 Overview

2 Tokyo example

3 GAG urine data

4 Close look at GAG urine example

- GAG reference limits

5 Details

GAG urine detailed

```
## prepare the data to predict new observations
gag.dat <- list(lgag=c(log(GAGurine$GAG), NA, NA),
               age=c(GAGurine$Age, 18, 19)) ## extrapolate !!!

## two models for the latent effect
gag.rw2.form <- lgag ~ f(inla.group(age, 100), model='rw2',
                        scale.model=TRUE, hyper=list(theta=list(prior='pc.prec', para
gag.sigm.form <- lgag ~ f(age, model='sigm')
## ask for more results
gag.rw2 <- inla(gag.rw2.form, data=gag.dat,
               control.predictor=list(compute=TRUE), ### to predict
               control.compute=list(config=TRUE, ## store configurations
                                   dic=TRUE, cpo=TRUE, po=TRUE), ## to compare
               control.inla=list(strategy='laplace')) ### best strategy
```

GAG model results

```
c(dic1=gag.rw2$dic$dic, lcpo1=-sum(log(gag.rw2$cpo$cpo), na.rm=TRUE),  
  dic2=gag.sigm$dic$dic, lcpo2=-sum(log(gag.sigm$cpo$cpo), na.rm=TRUE))
```

```
##          dic1          lcpo1          dic2          lcpo2  
## 136.83245   71.59637 175.26973   88.44898
```

```
round(gag.rw2$summary.hy[, 1:2], 4)
```

```
##                                     mean      sd  
## Precision for the Gaussian observations 11.4984  0.9310  
## Precision for inla.group(age, 100)      15.2572 12.4783
```

```
round(gag.sigm$summary.hy[, 1:2], 4)
```

```
##                                     mean      sd  
## Precision for the Gaussian observations 10.0989 0.8266  
## SIGM beta for age                     -1.8636 0.1119  
## SIGM halflife for age                   4.0201 0.4520  
## SIGM shape for age                     1.2921 0.1229
```

Prediction-extrapolation

```
cbind(true=c(GAGurine$GAG[314], NA, NA), # with 314
      rw2=exp(gag.rw2$summary.fitted.val[314:316, 1]),
      sigm=exp(gag.sigm$summary.fitted.val[314:316, 1]))
```

```
##      true      rw2      sigm
## [1,]  9.3 2.954101 4.803916
## [2,]   NA 2.877849 4.780320
## [3,]   NA 2.659282 4.714983
```

Error in the linear predictor scale

```
rbind(rw2=gag.rw2$summary.fitted.val[314:316, 2],
      sigm=gag.sigm$summary.fitted.val[314:316, 2])
```

```
##           [,1]      [,2]      [,3]
## rw2 0.13067358 0.16259568 0.28892200
## sigm 0.04275026 0.04332279 0.04499502
```

The variance

- compute the distribution for $\sigma|y$ (std.dev.)

```
tau.marginal <- gag.rw2$marginals.hyperpar[[1]]  
s.marginal <- inla.tmarginal(function(x) sqrt(1/x), tau.marginal)
```

- summarize a marginal

```
inla.zmarginal(s.marginal)  
  
## Mean          0.295627  
## Stdev         0.0119119  
## Quantile 0.025 0.273161  
## Quantile 0.25  0.287379  
## Quantile 0.5   0.295262  
## Quantile 0.75  0.303493  
## Quantile 0.975 0.31992
```

- other functions to work with marginals

```
apropos('marginal')  
  
## [1] "inla.dmarginal" "inla.emarginal" "inla.hpdmarginal"  
## [4] "inla.mmarginal" "inla.pmarginal" "inla.qmarginal"  
## [7] "inla.rmarginal" "inla.smarginal" "inla.tmarginal"  
## [10] "inla.zmarginal" "s.marginal"     "tau.marginal"
```

GAG limits: The problem

- Define limits (2.5%, 95.5%) for GAG for each age
- The statistical model is for the mean: $E(y)$
- Additionally we have the *error*: $GAG = E(GAG) + \text{error}$

GAG limits: The problem

- Define limits (2.5%, 95.5%) for GAG for each age
- The statistical model is for the mean: $E(y)$
- Additionally we have the *error*: $GAG = E(GAG) + \text{error}$
- Easy way solution:
 - sample from $E(y)$ and the error
 - compute the sum
 - define the interval

GAG limits: code

get samples from the posterior distribution

```
samples <- inla.posterior.sample(n=10000, res=gag.rw2)
```


for each sample, sample the error and compute the sum

```
psam <- sapply(samples, function(x) {  
  s.d <- sqrt(1/x$hyperpar[1])  
  mu <- x$latent[1:316,1]  
  err <- rnorm(316, 0, s.d)  
  return(exp(mu+err)) ## in the GAG original scale  
})  
q.sam <- apply(psam, 1, quantile,  
  c(.025, 0.05, 0.5, 0.95, .975))
```

```

par(mar=c(2.5, 2.5, 0, 0), mgp=c(1.5, 0.5, 0))
plot(gag.dat$age, c(GAGurine$GAG,NA,NA), pch=19,
     col=gray(.5), xlab='Age (years)',
     ylab='Childrens GAG concentration')
for (j in 1:5)
  lines(gag.dat$age, q.sam[j, ], lty=c(3,2,1,2,3)[j], lwd=2)

```



Visualize GAG limits

Outline

- 1 Overview
- 2 Tokyo example
- 3 GAG urine data
- 4 Close look at GAG urine example
- 5 Details

Results, post-processing and sampling

- Accessing results:
 - `summary(result)`
 - `plot(result)`
- post-processing
 - `result2 = inla.hyperpar(result)`
 - `result2 = inla.cpo(result)`
 - `result2 = inla.rerun(result)`
- sampling from posterior joint distribution
 - `samples = inla.hyperpar.sample(n=1000, result)`
 - `samples = inla.posterior.sample(n=1000, result)`

The interpretation of NA

INLA uses NA differently than other packages

- NA in the response means no likelihood contribution, i.e. response is unobserved
- NA in a fixed effect means no contribution to the linear predictor, i.e. the covariate is set equal to zero
- NA in a random effect $f(\dots)$ means no contribution to the linear predictor

Changing the prior: Internal scale

- Hyperparameters are represented internally with more well-behaved transformations, e.g. correlation ρ and precision τ are internally represented as

$$\theta_1 = \log(\tau)$$

$$\theta_2 = \log\left(\frac{1 + \rho}{1 - \rho}\right)$$

- The prior must be set on the parameter in **internal scale**
- Initial values for the mode-search must be set in **internal scale**
- The functions `to.theta` and `from.theta` can be used to map back and forth.

Changing the prior: Code

```
## define the prior
hyper = list(prec = list(prior = "loggamma",
                          param = c(1, 0.1),
                          initial = 4,
                          fixed = FALSE))

## insert it into the formula
formula = y ~ f(idx, model = "iid", hyper = hyper) + ...
```

Changing the prior: Default options

```
## Default options can be seen with  
inla.models()$latent$id$hyper
```

```
theta  
  name "log precision"  
short.name "prec"  
  prior "loggamma"  
  param c(1e+00, 5e-05)  
initial 4  
  fixed FALSE  
to.theta function(x){log(x)}  
from.theta function(x){exp(x)}
```

Available priors

```
names(inla.models())$prior)
```

```
## [1] "normal" "gaussian"
## [3] "wishart1d" "wishart2d"
## [5] "wishart3d" "wishart4d"
## [7] "wishart5d" "loggamma"
## [9] "minuslogsqrtruncnormal" "logtnormal"
## [11] "logtgaussian" "flat"
## [13] "logflat" "logiflat"
## [15] "mvnorm" "pc.ar"
## [17] "none" "invalid"
## [19] "betacorrelation" "logitbeta"
## [21] "pc.prec" "pc.dof"
## [23] "pc.cor0" "pc.cor1"
## [25] "pc.fgnh" "pc.spde.GA"
## [27] "pc.matern" "pc.range"
## [29] "pc.gamma" "pc.mgamma"
## [31] "pc.gammacount" "pc"
## [33] "ref.ar" "jeffreystdf"
## [35] "expression:" "table:"
```

Available likelihoods

```
names(inla.models())$likelihood
```

```
## [1] "poisson" "qpoisson"
## [3] "cenpoisson" "gpoisson"
## [5] "binomial" "testbinomial1"
## [7] "gamma" "gammacount"
## [9] "qkumar" "qloglogistic"
## [11] "beta" "betabinomial"
## [13] "cbinomial" "nbinomial"
## [15] "simplex" "gaussian"
## [17] "normal" "circularnormal"
## [19] "wrappedcauchy" "iidgamma"
## [21] "iidlogitbeta" "loggammafrailty"
## [23] "logistic" "skewnormal"
## [25] "sn" "sn2"
## [27] "gev" "lognormal"
## [29] "lognormalsurv" "exponential"
## [31] "exponentialsurv" "coxph"
## [33] "weibull" "weibullsurv"
## [35] "loglogistic" "weibullcure"
## [37] "stochvol" "stochvolt"
## [39] "stochvolnig" "zeroinflatedpoisson0"
## [41] "zeroinflatedpoisson1" "zeroinflatedpoisson2"
```


Available latent models

```
names(inla.models())$latent)
```

## [1]	"linear"	"iid"	"mec"	"meb"
## [5]	"rgeneric"	"rw1"	"rw2"	"crw2"
## [9]	"seasonal"	"besag"	"besag2"	"bym"
## [13]	"bym2"	"besagproper"	"besagproper2"	"fgn"
## [17]	"ar1"	"ar"	"ou"	"generic"
## [21]	"generic0"	"generic1"	"generic2"	"generic3"
## [25]	"spde"	"spde2"	"spde3"	"iid1d"
## [29]	"iid2d"	"iid3d"	"iid4d"	"iid5d"
## [33]	"2diid"	"z"	"rw2d"	"rw2diid"
## [37]	"slm"	"matern2d"	"copy"	"clinear"
## [41]	"sigm"	"revsigm"	"log1exp"	"logdist"